Object Detection Using CNNs

C.-C. Jay Kuo
University of Southern California

History in Object Detection

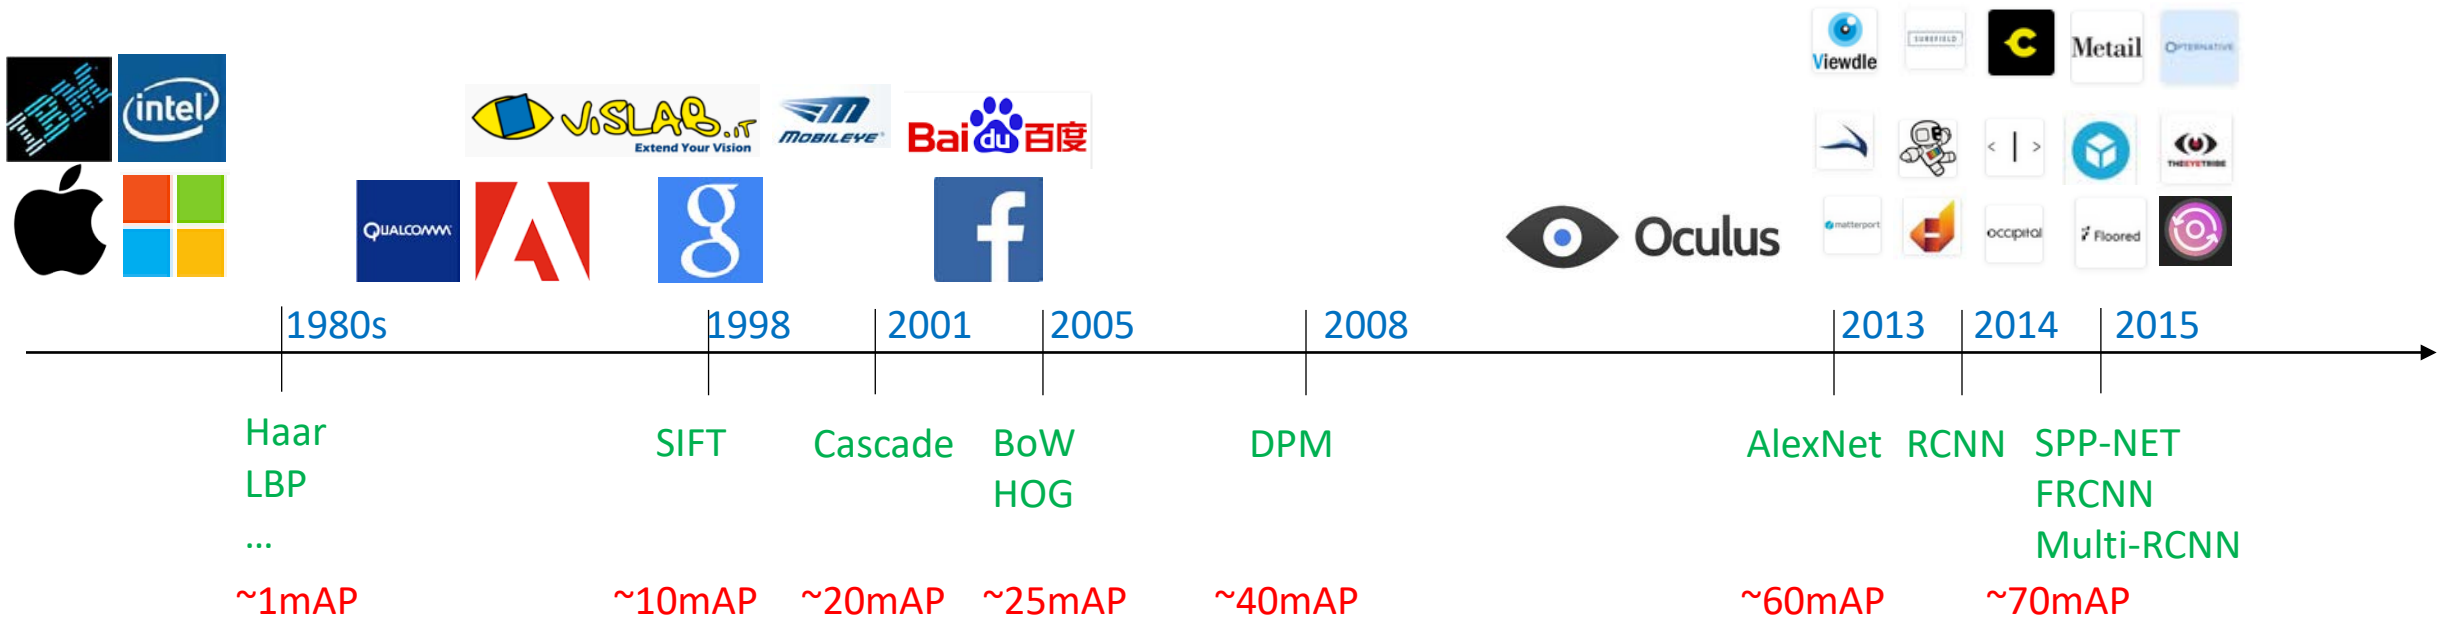

mAP: Mean Average Precision

Object Detection with DNN

2013

AlexNet

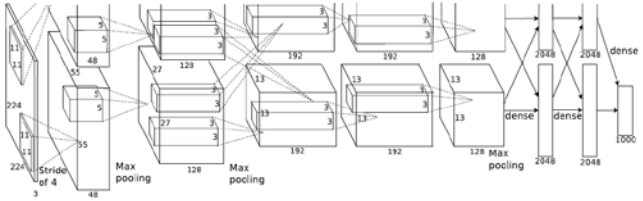

2014

RCNN

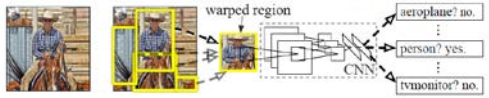

Caffe

DenseNet

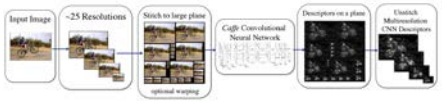

Part-Based R-CNN

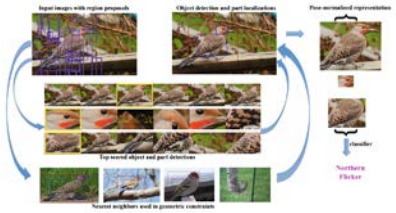

2015

SPP-Net

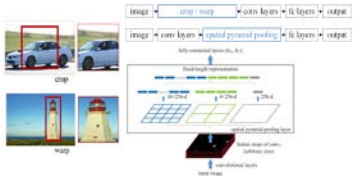

Fast R-CNN

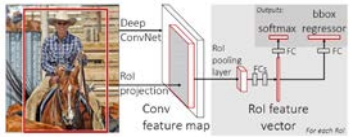

Faster R-CNN

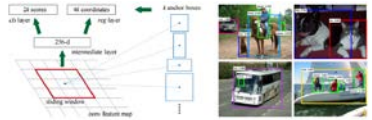

Alex Net (2013)

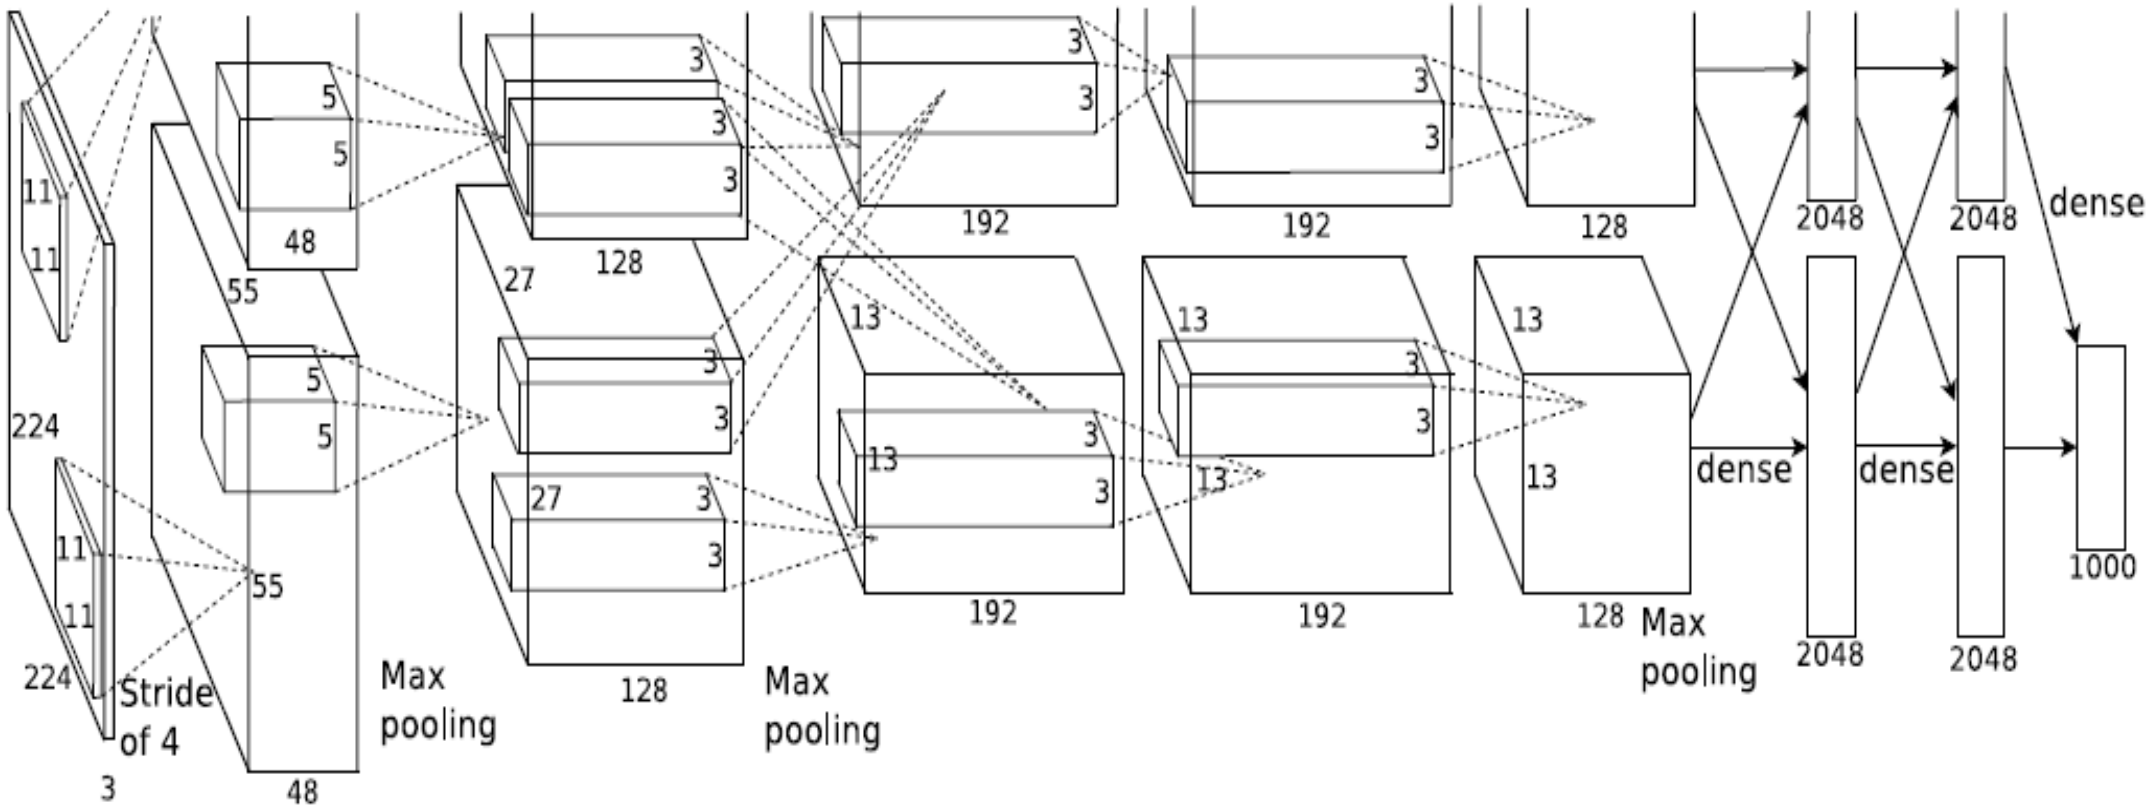

A. Krizhevsky, I. Sutskever, and G. E. Hinton. Imagenet classification with deep convolutional neural networks.

Region Proposal with CNN (RCNN)

Constraint Parametric Min-Cut

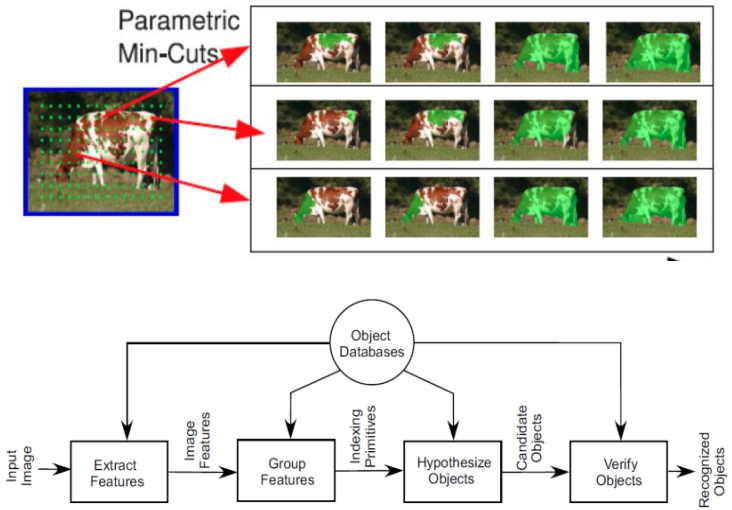

J. Carreira and C.sminchisescu. CPMC Automatic object segmentation using constrained parametric min-cust. In CVPR, 2012

Selective Search

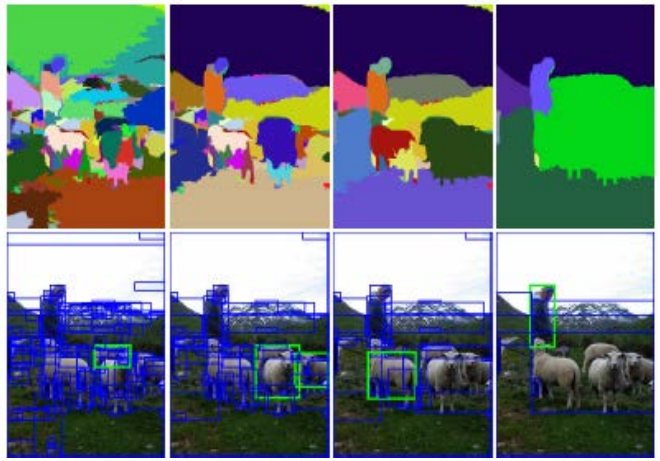

J.R.R Uijlings, K.E.A. Van De Sande, T. Gevers, and A.W.M. Smeulders, Selective Search for Object Recognition. In IJCV 2013

EdgeBox

C. L. Zitnick and P. Dollar, “Edge boxes: Locating object ´ proposals from edges,” in ECCV, 2014.

RCNN Flowchart

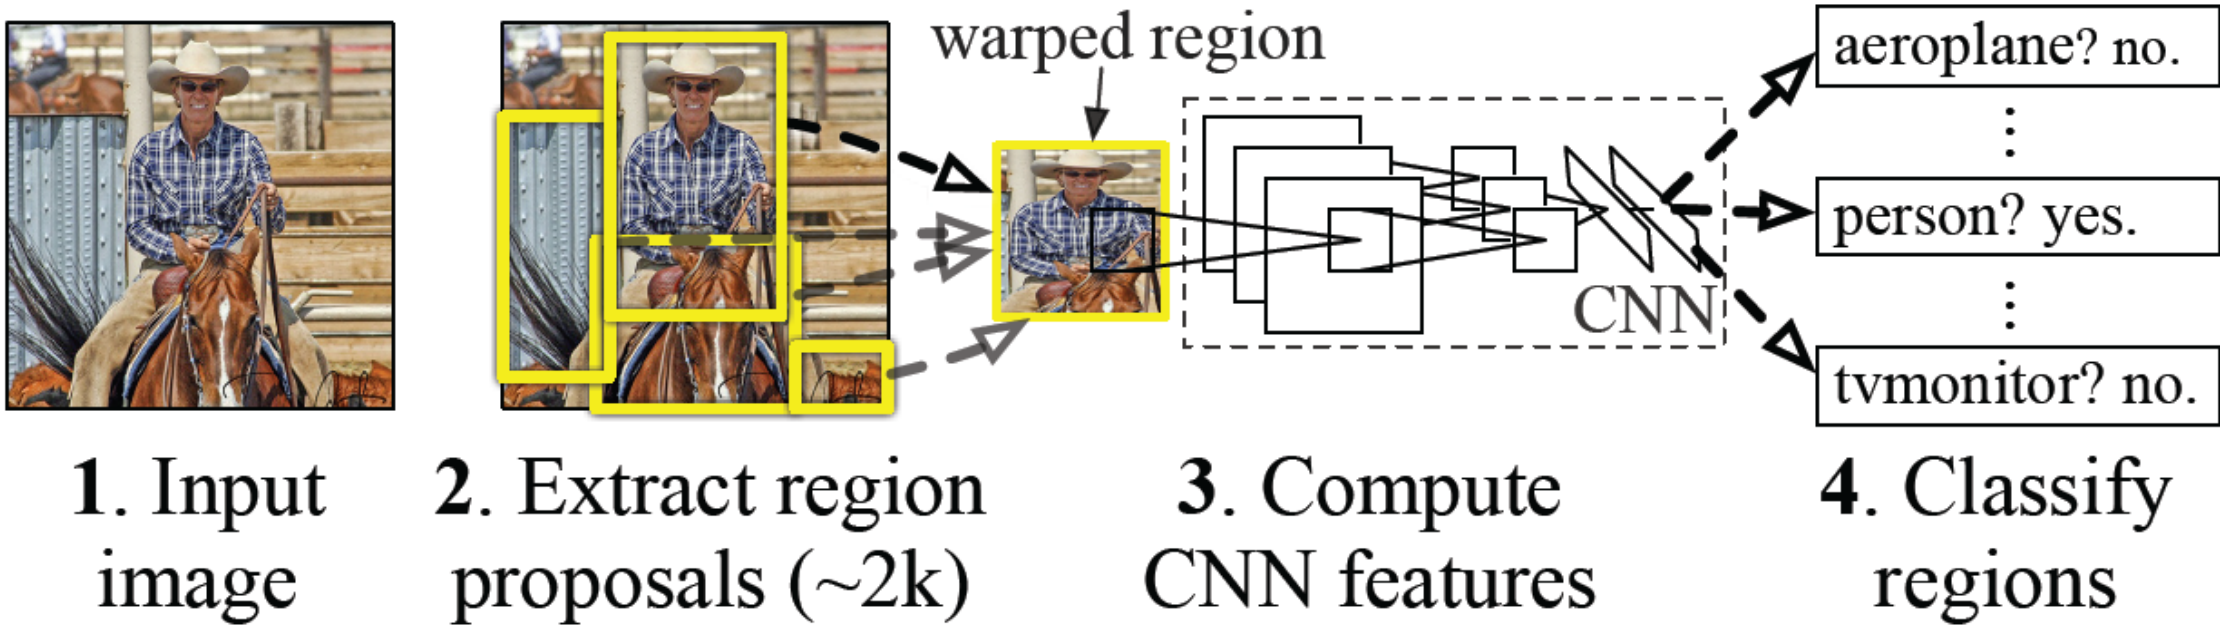

Ross Girshick, Jeff Donahue, Trevor Darrell, Jitendra Malik, Rich Feature Hierachies for Accurate Object Detection and Semantic Segmentation

Caffe: Public Domain Software

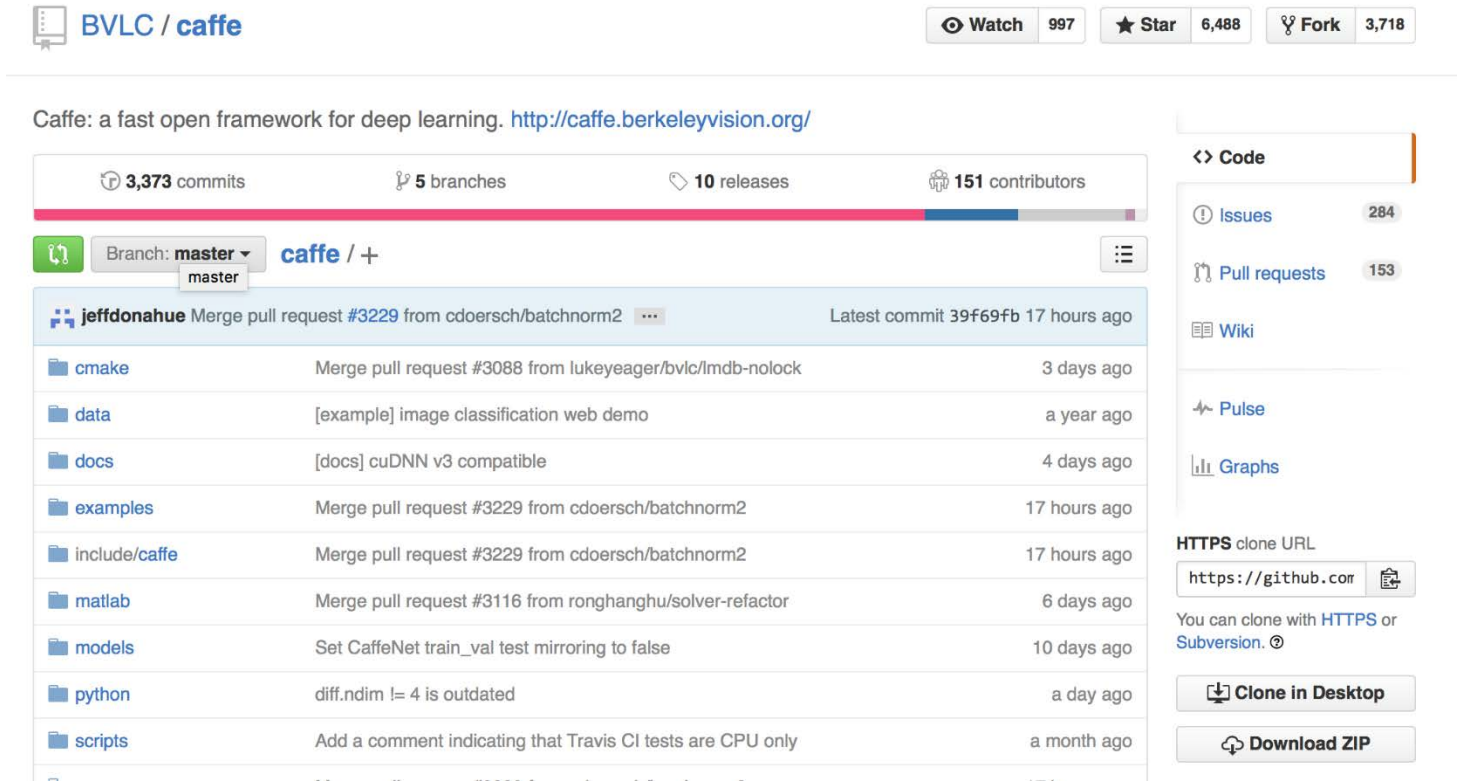

R. Girshick, J. Donahue, T. Darrell, and J. Malik. Rich feature hierarchies for accurate object detection and semantic segmentation.

Dense Net

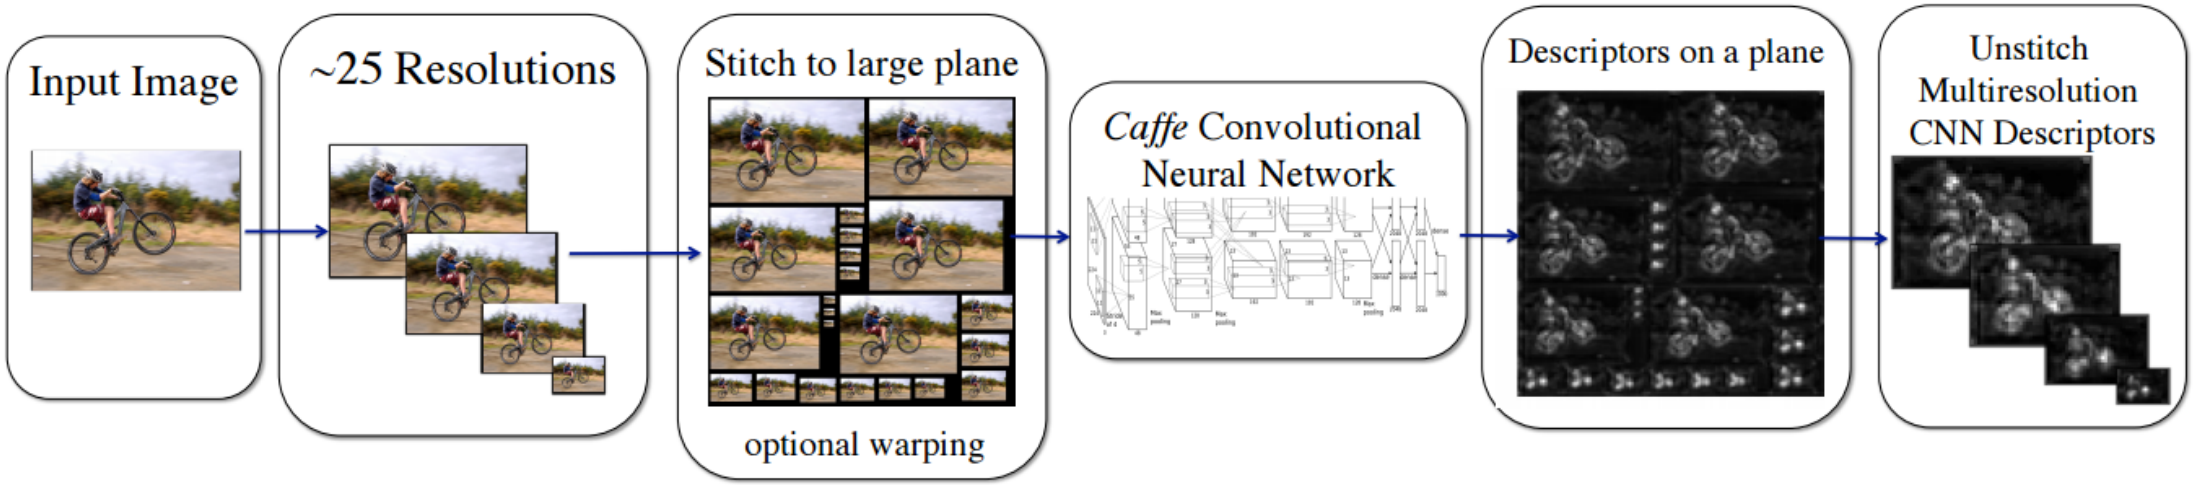

F.Iandola, M. Moskewicz, S. Karayev, R. Girshick, T. Darrell, K. Keutzer, DenseNet: Implementing Efficient ConvNet Descriptor Pyramids

Part-based CNN

- Fine-grained Detection

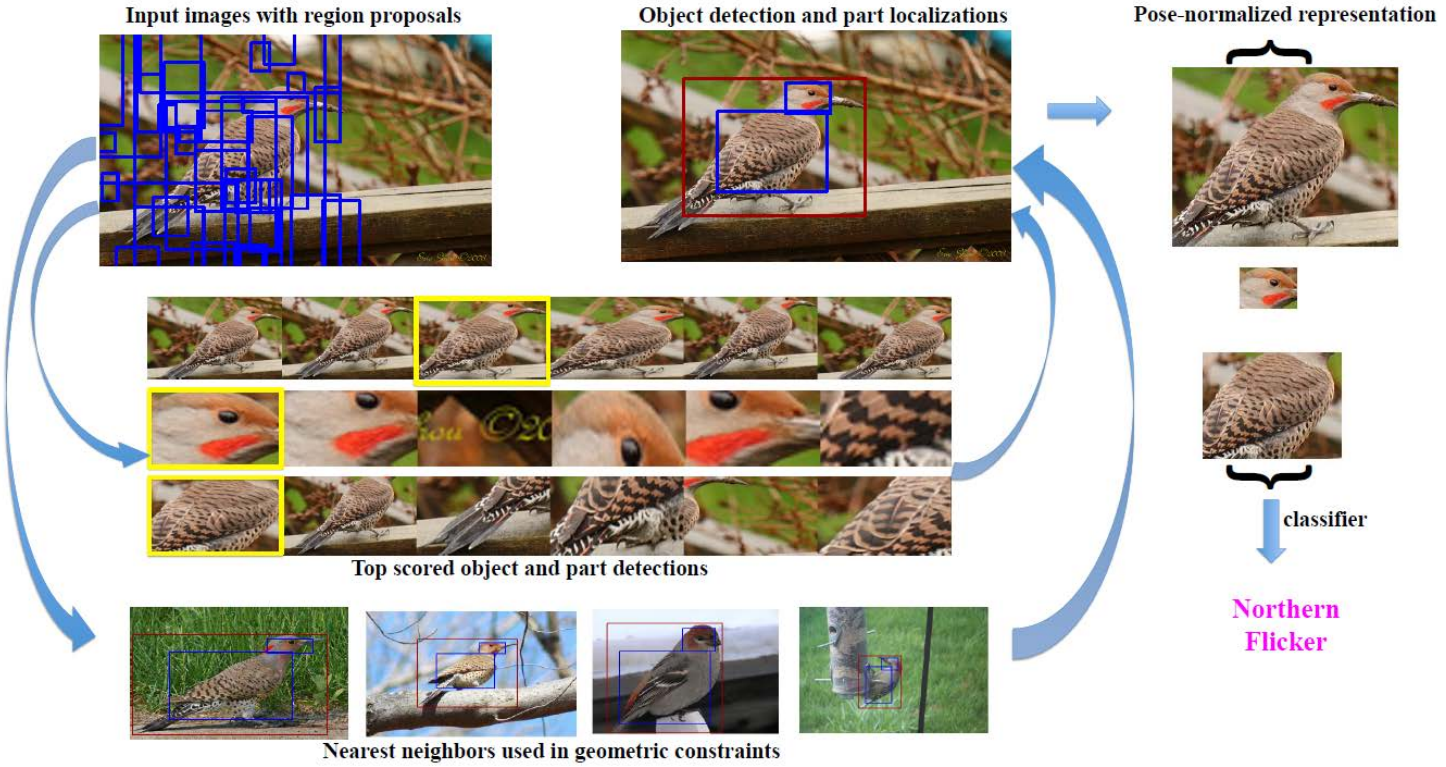

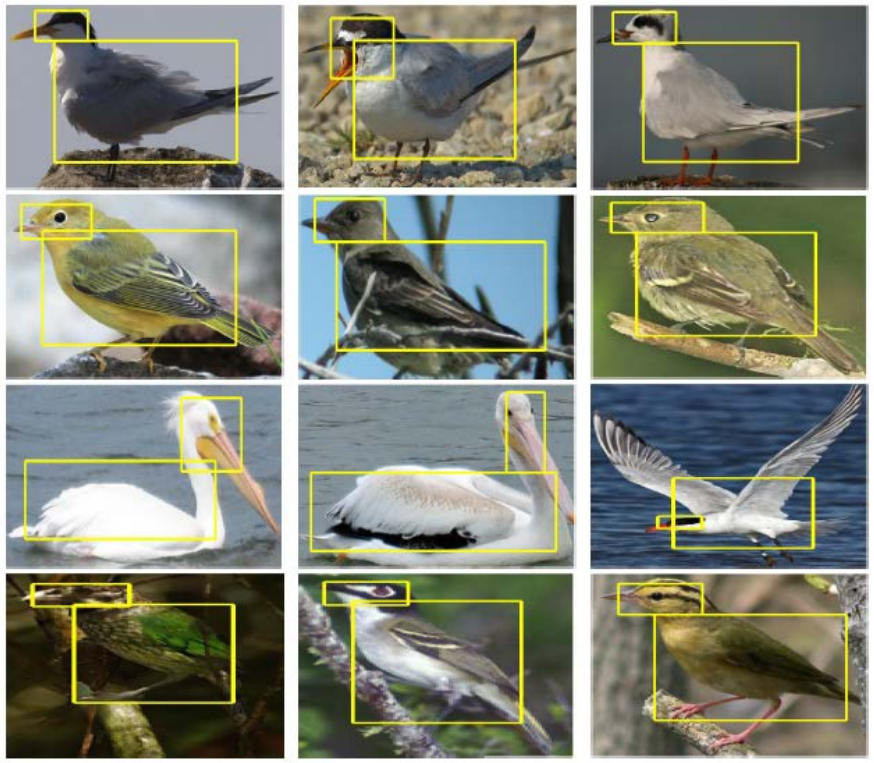

Multi-region CNN

- Multi-region & semantic segmentation-aware CNN model:

(a) Original box

(b) Half left

(c) Half right

(d) Half up

(e) Half bottom

(f) Central Region

(g) Central Region

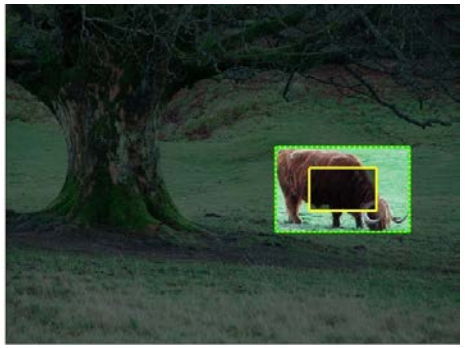

(h) Border Region

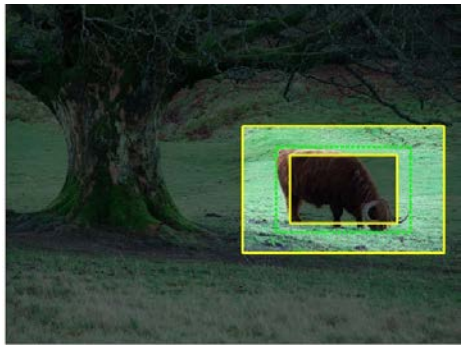

(i) Border Region

(j) Context. Region

Fast RCNN

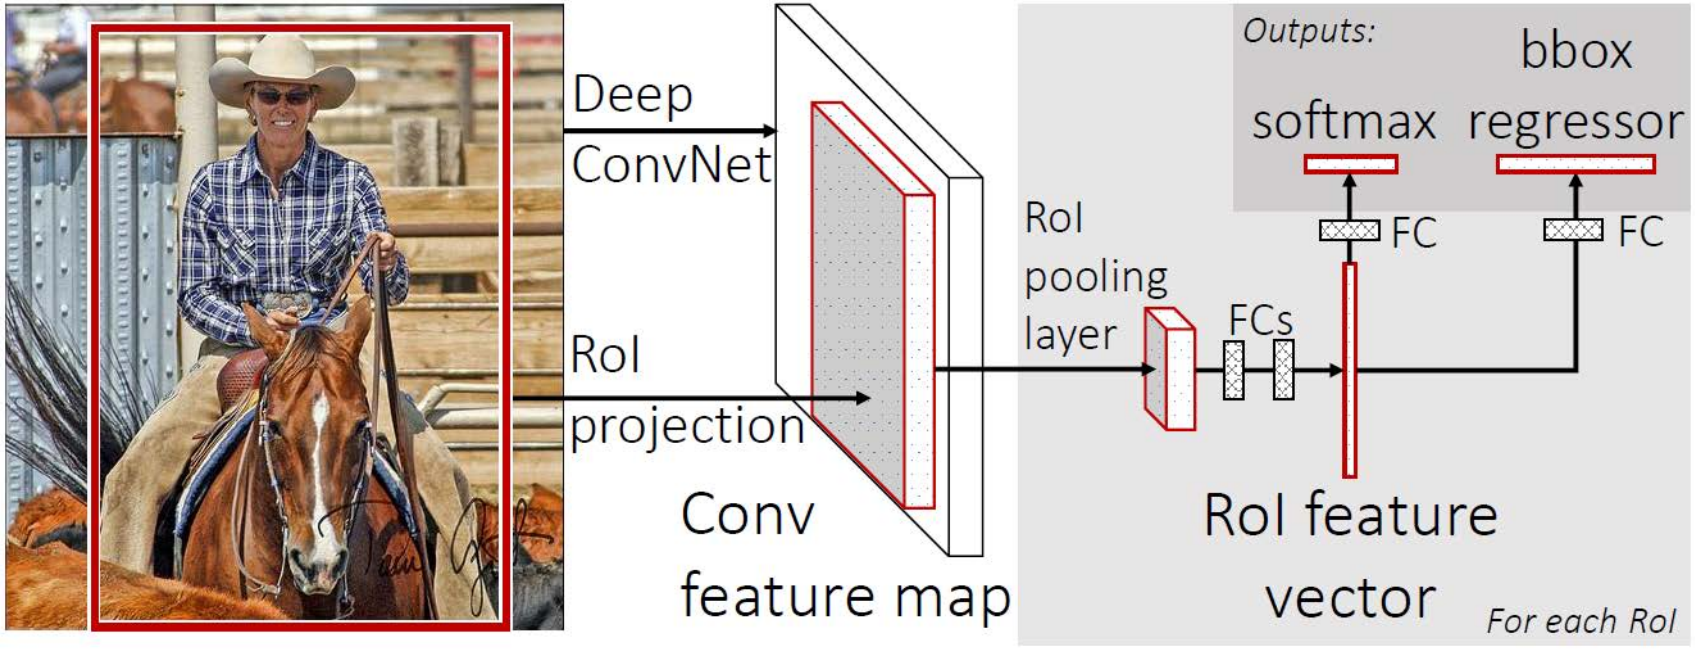
Deep
ConvNet
RoI
projection
Conv
feature map
RoI
pooling
layer
FCs
Outputs:
softmax
bbox
regressor
FC
FC
RoI feature
vector
For each RoI

Faster RCNN

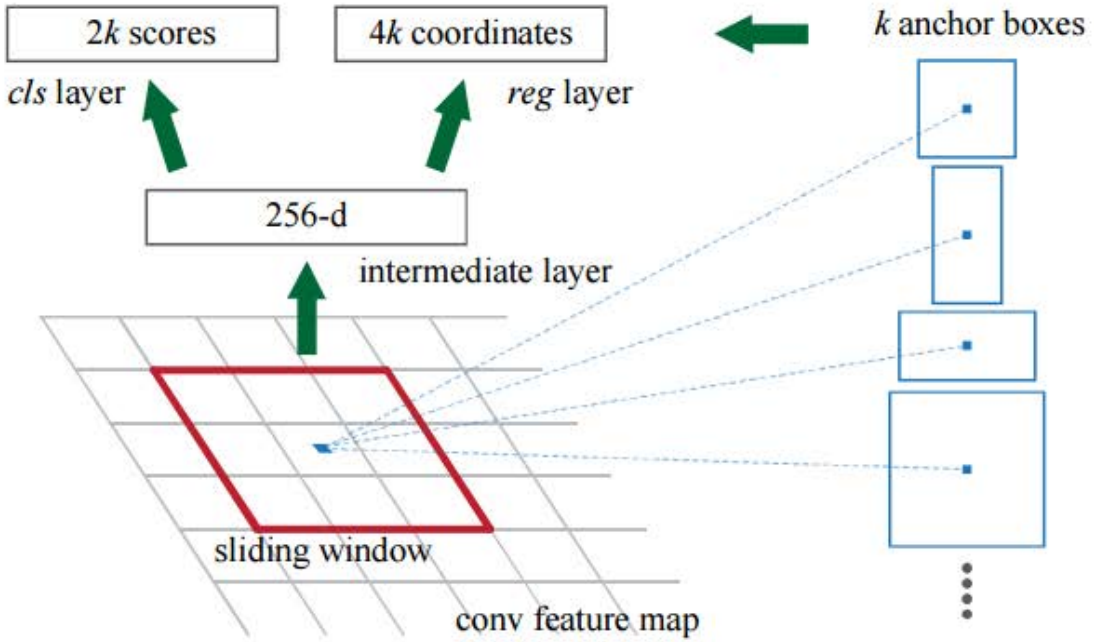

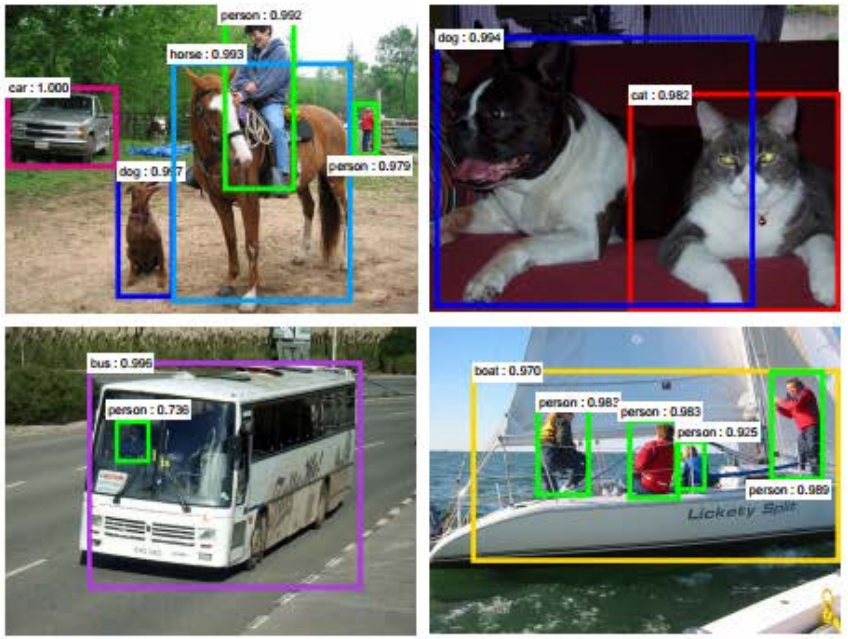

S. Ren, K. He, R. Girshick, and J. Sun. Faster r-cnn: Towards real-time object detection with region proposal networks.

YOLO

- More improved CNN based Object Detector:

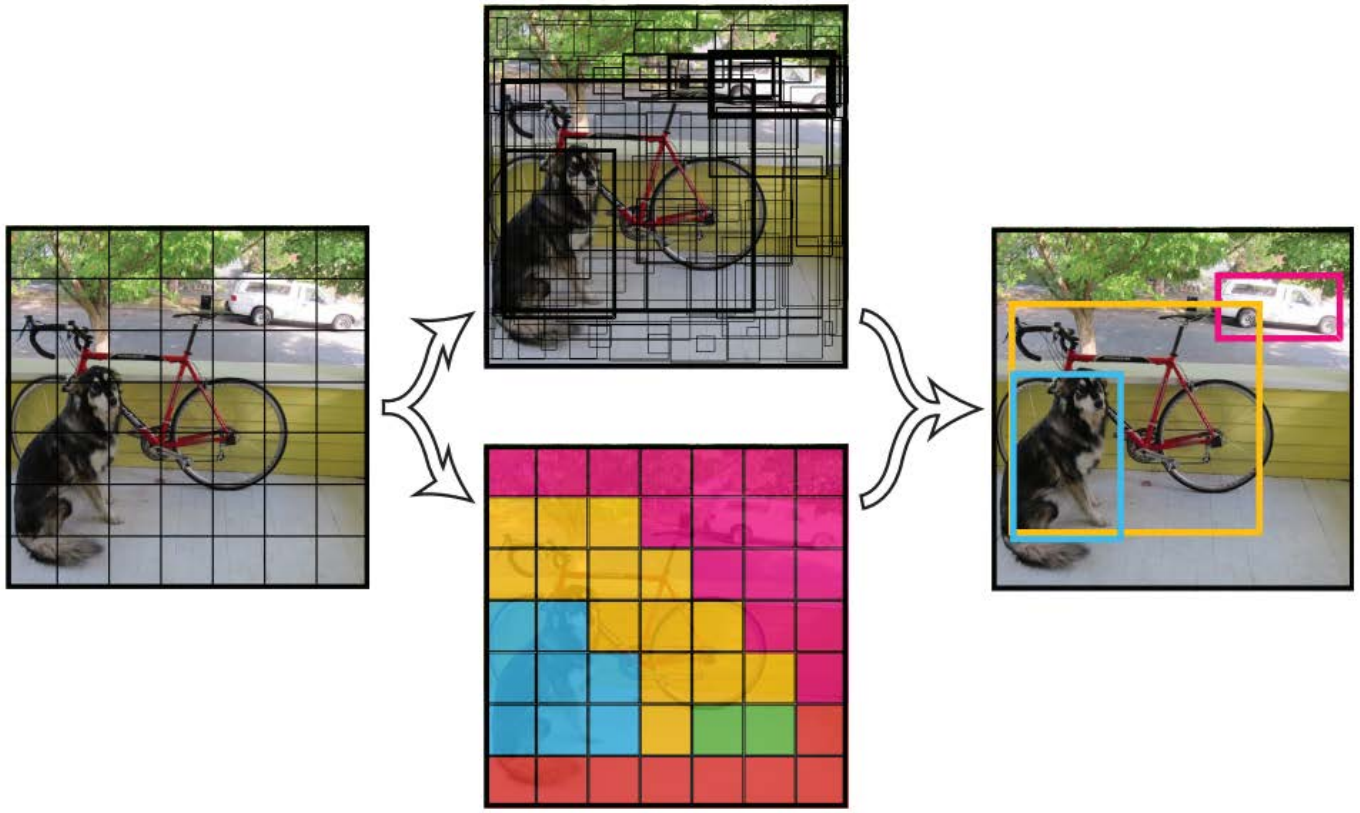

Joseph Redmon, Santosh Divvala, Ross Girshick, Ali Farhadi, You Only Look Once: Unified, Real-Time Object Detection

Performance Comparison

- State-of-the-art methods comparison:

Method	ROI needed?	Regression?	Accuracy	Speed
RCNN	Y	N	62.4	0.01 fps
Fast-RCNN	Y	Y	68.4	0.01 fps
Faster-RCNN	N	Y	70.4	3 fps
YOLO	N	Y	57.9	45 fps

Comments

- Alex Net was the first one proposed for object detection
- Recent publications have focused on two aspects:
  - Fine-tuning for better region proposals, which serve as a critical pre-processing step for DNN
  - Seeking more applications
- Both hardware and software are accessible to the public
- However, no labeled training data are accessible to the public
- It can be implemented in the cloud platform, but not in the embedded system platform

Visualization of Deep Features

Feature Visualization

- De-conv Network:

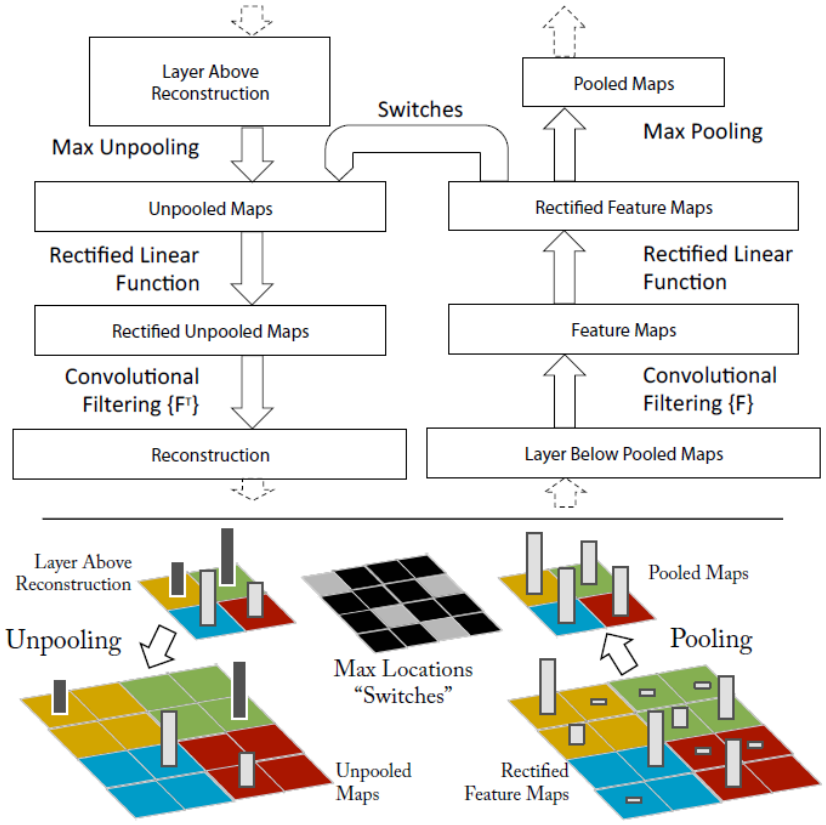

Features at Lower Layers

- Conv1 Filter Response (Gabor-like filters)

Figure Credit: Stanford CS231n Course Instruction

Reconstruction at Lower Layers

- Extracting low level features (oriented edges/contours)

http://googleresearch.blogspot.ch/2015/06/inceptionism-going-deeper-into-neural.html

Reconstruction at Higher Layers

- Random Noise Input

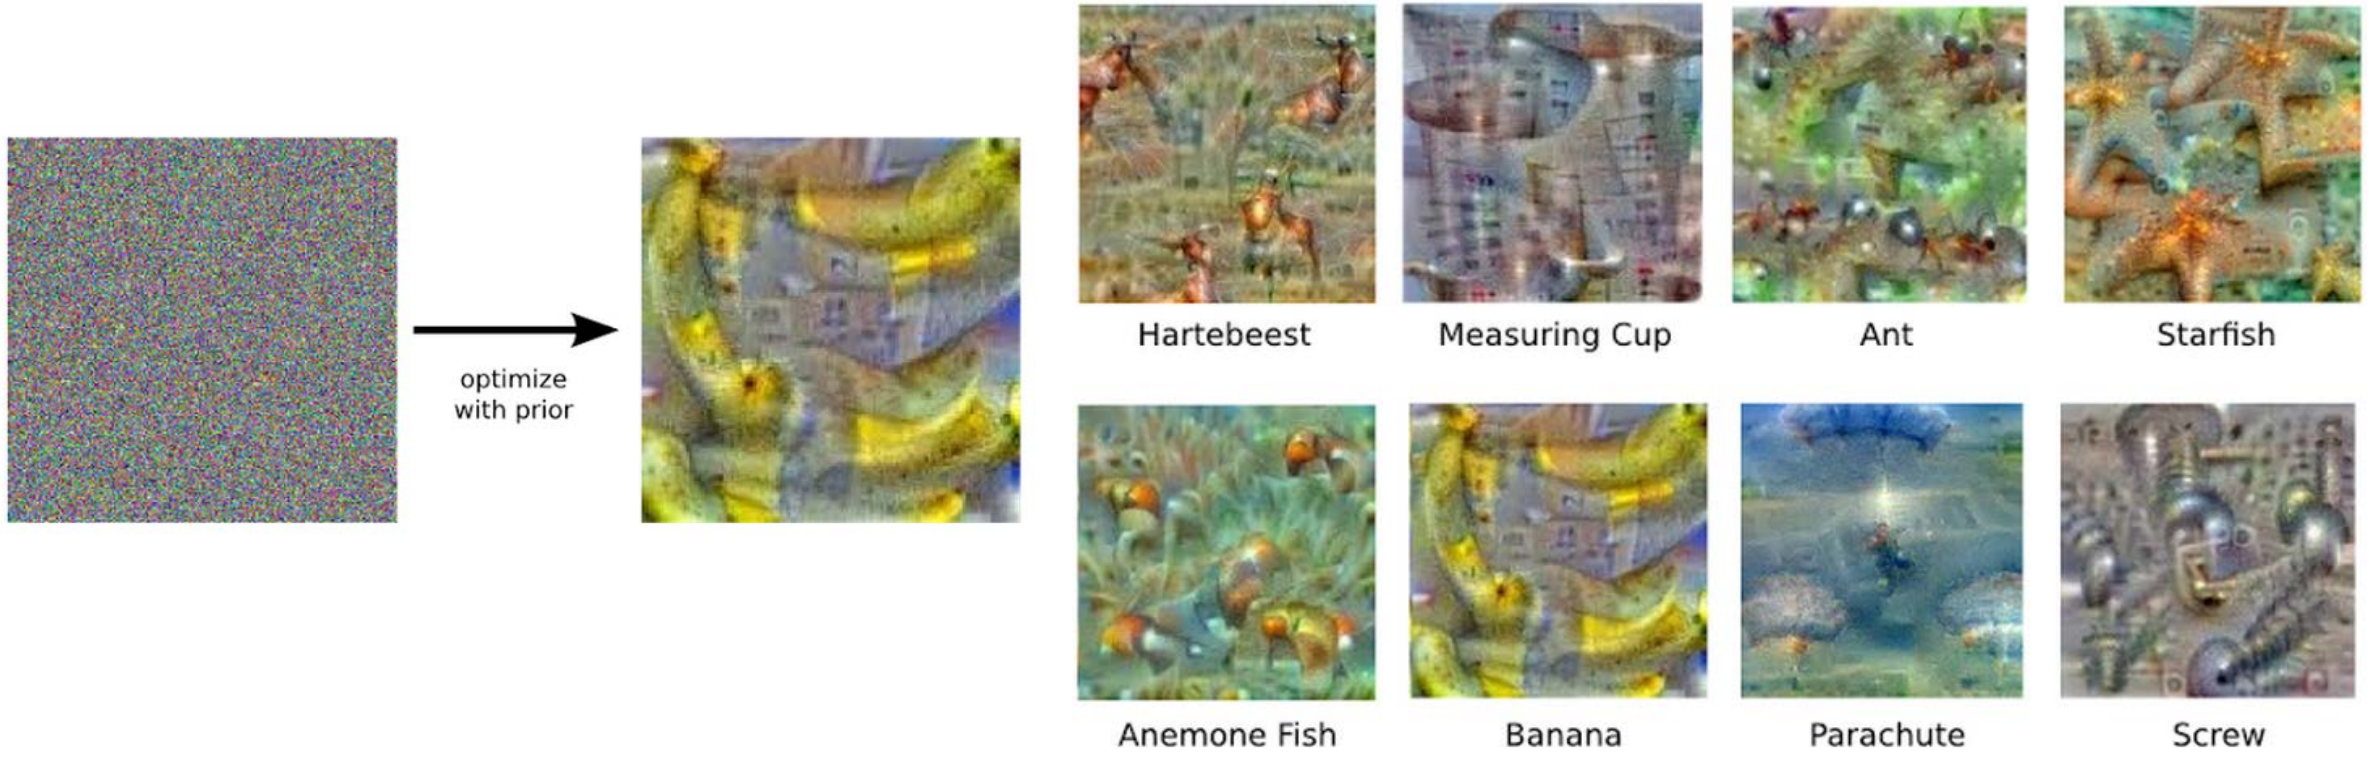

http://googleresearch.blogspot.ch/2015/06/inceptionism-going-deeper-into-neural.html

Deep Features at Conv5

- Cat Face Filter Example (Filter 111 in the Yos-CaffeNet Model):

Top 9 Input Activation Images

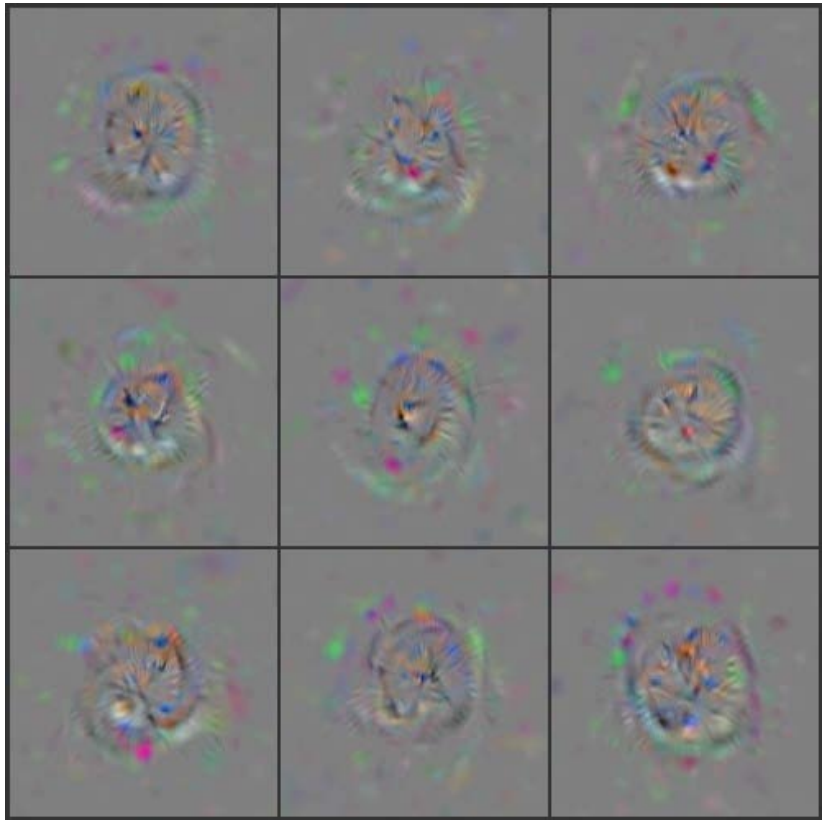

Max Reconstructed Input Activation

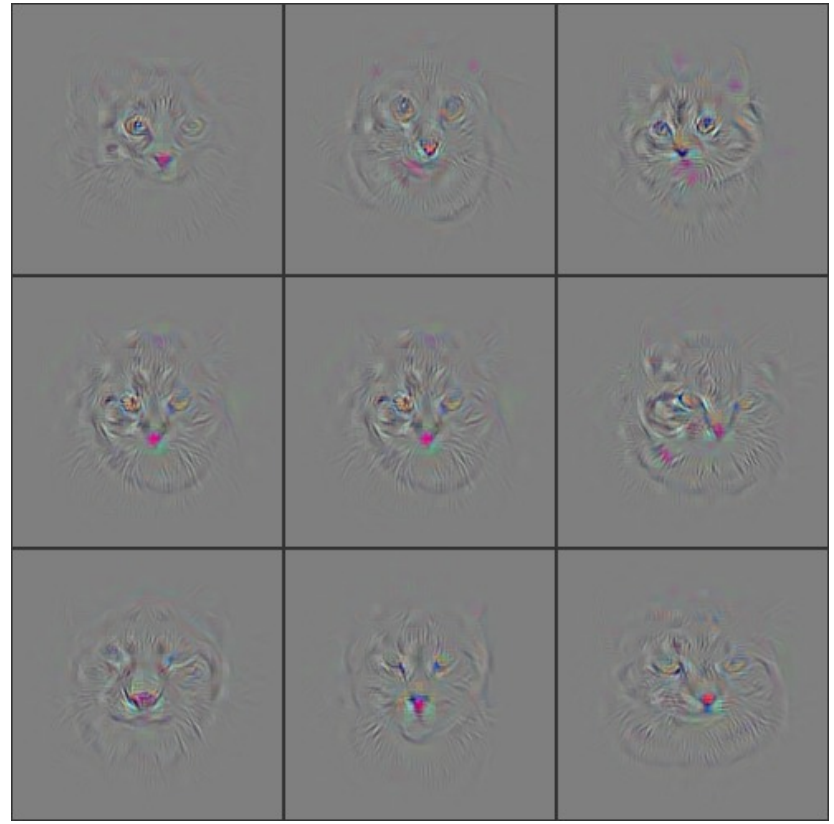

Deconv Image

Images are generated from Deep-VIS Toolbox

Comments

- Deep feature visualization provides helpful insight into the role played by various layers
  - Lower layers – subband-filtered images with large coverage
  - Higher layers – contour-dominating images with focused coverage
  - Three general trends
    - Contour formation
      - From surfaces to contours
    - Color reduction
      - From colorful to colorless
    - Background removal
      - From larger regions (with background) to focused regions (without background)
- The whole process can be viewed as a spatial (contour) feature binding process

Deep Features for Scene Recognition

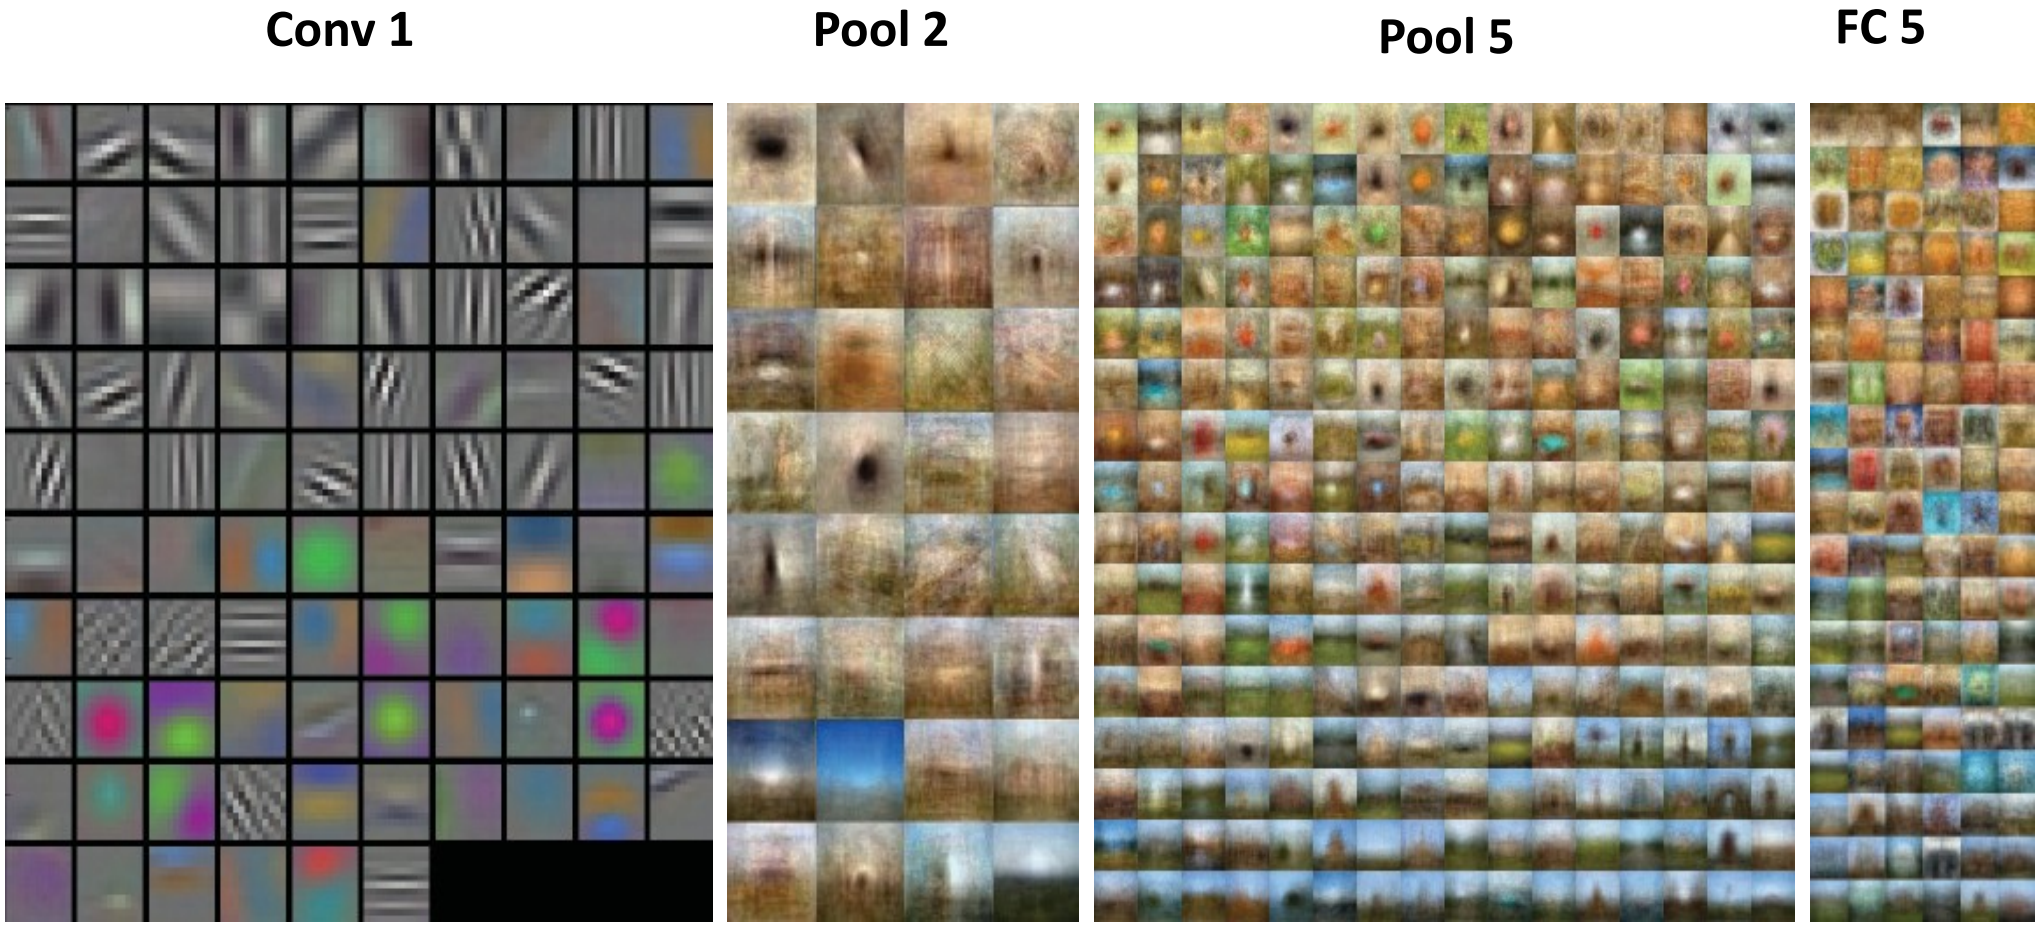

Zhou et al. “Learning deep features for scene recognition using places database”, NIPS 2014.

Gaussian Confusion Measure (GCM)

Conv5 Filter Response

- Gaussian Confusion Measure

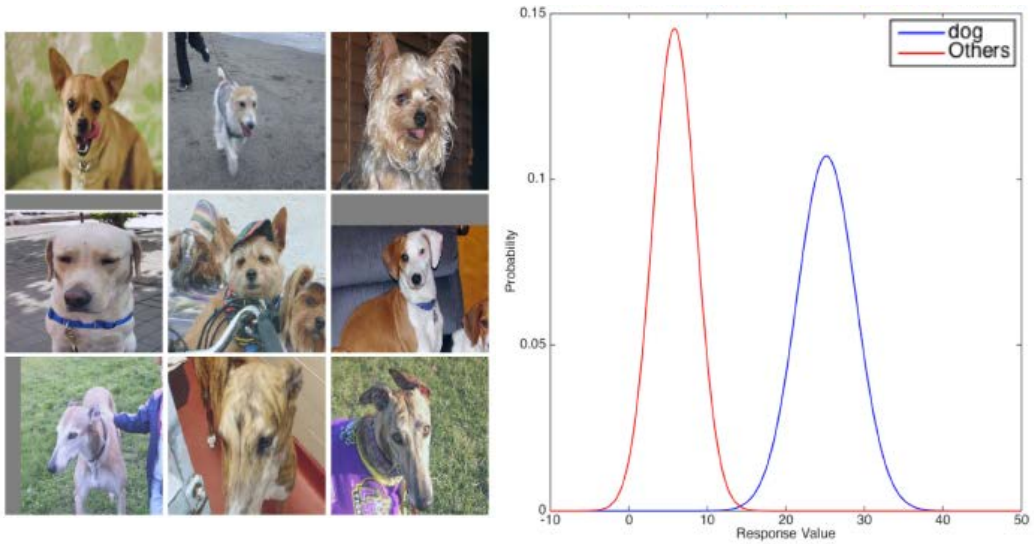

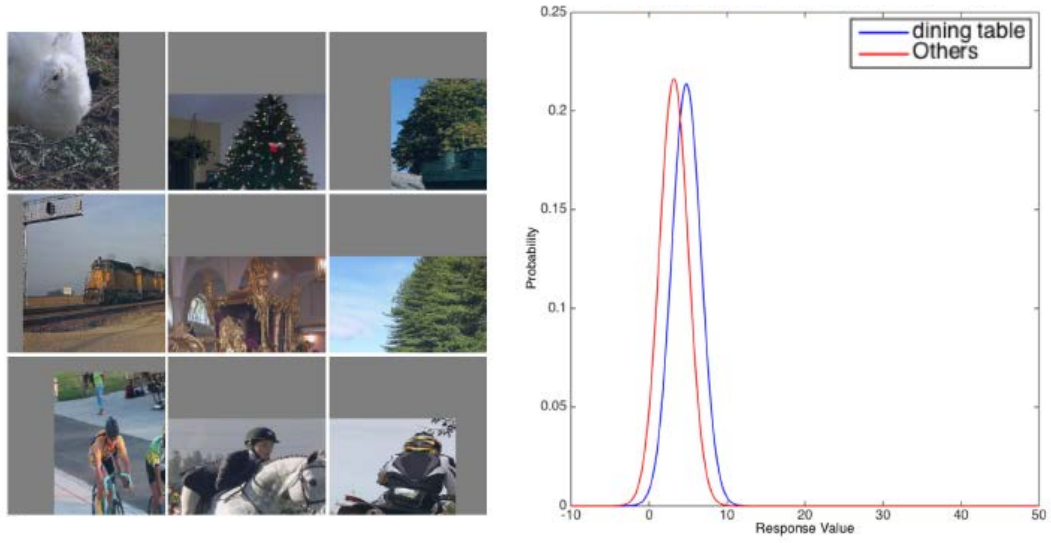

Detection Error and Mean of Conv5 Filter Responses

Responses w.r.t. human objects

-- Known as “Grandmother Cell (GMC)” like features for human objects

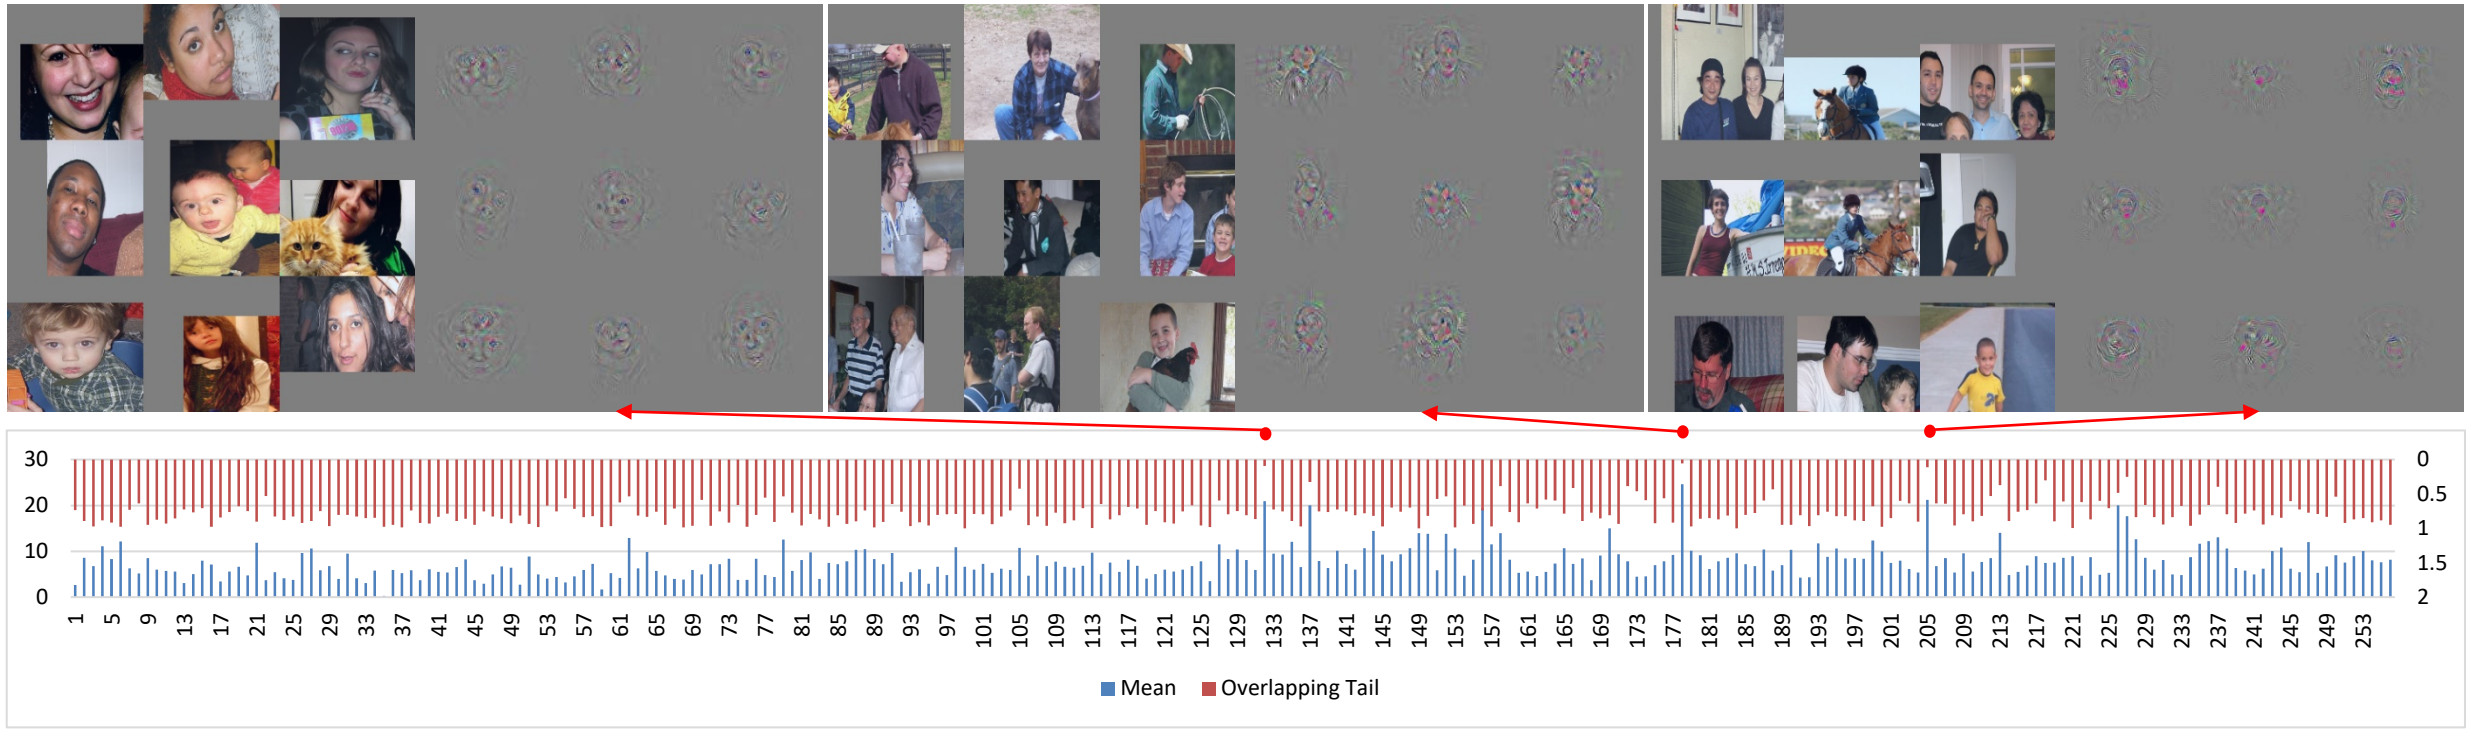

Optimal Conv5 Filters Against Multiple Object Classes (Good Examples)

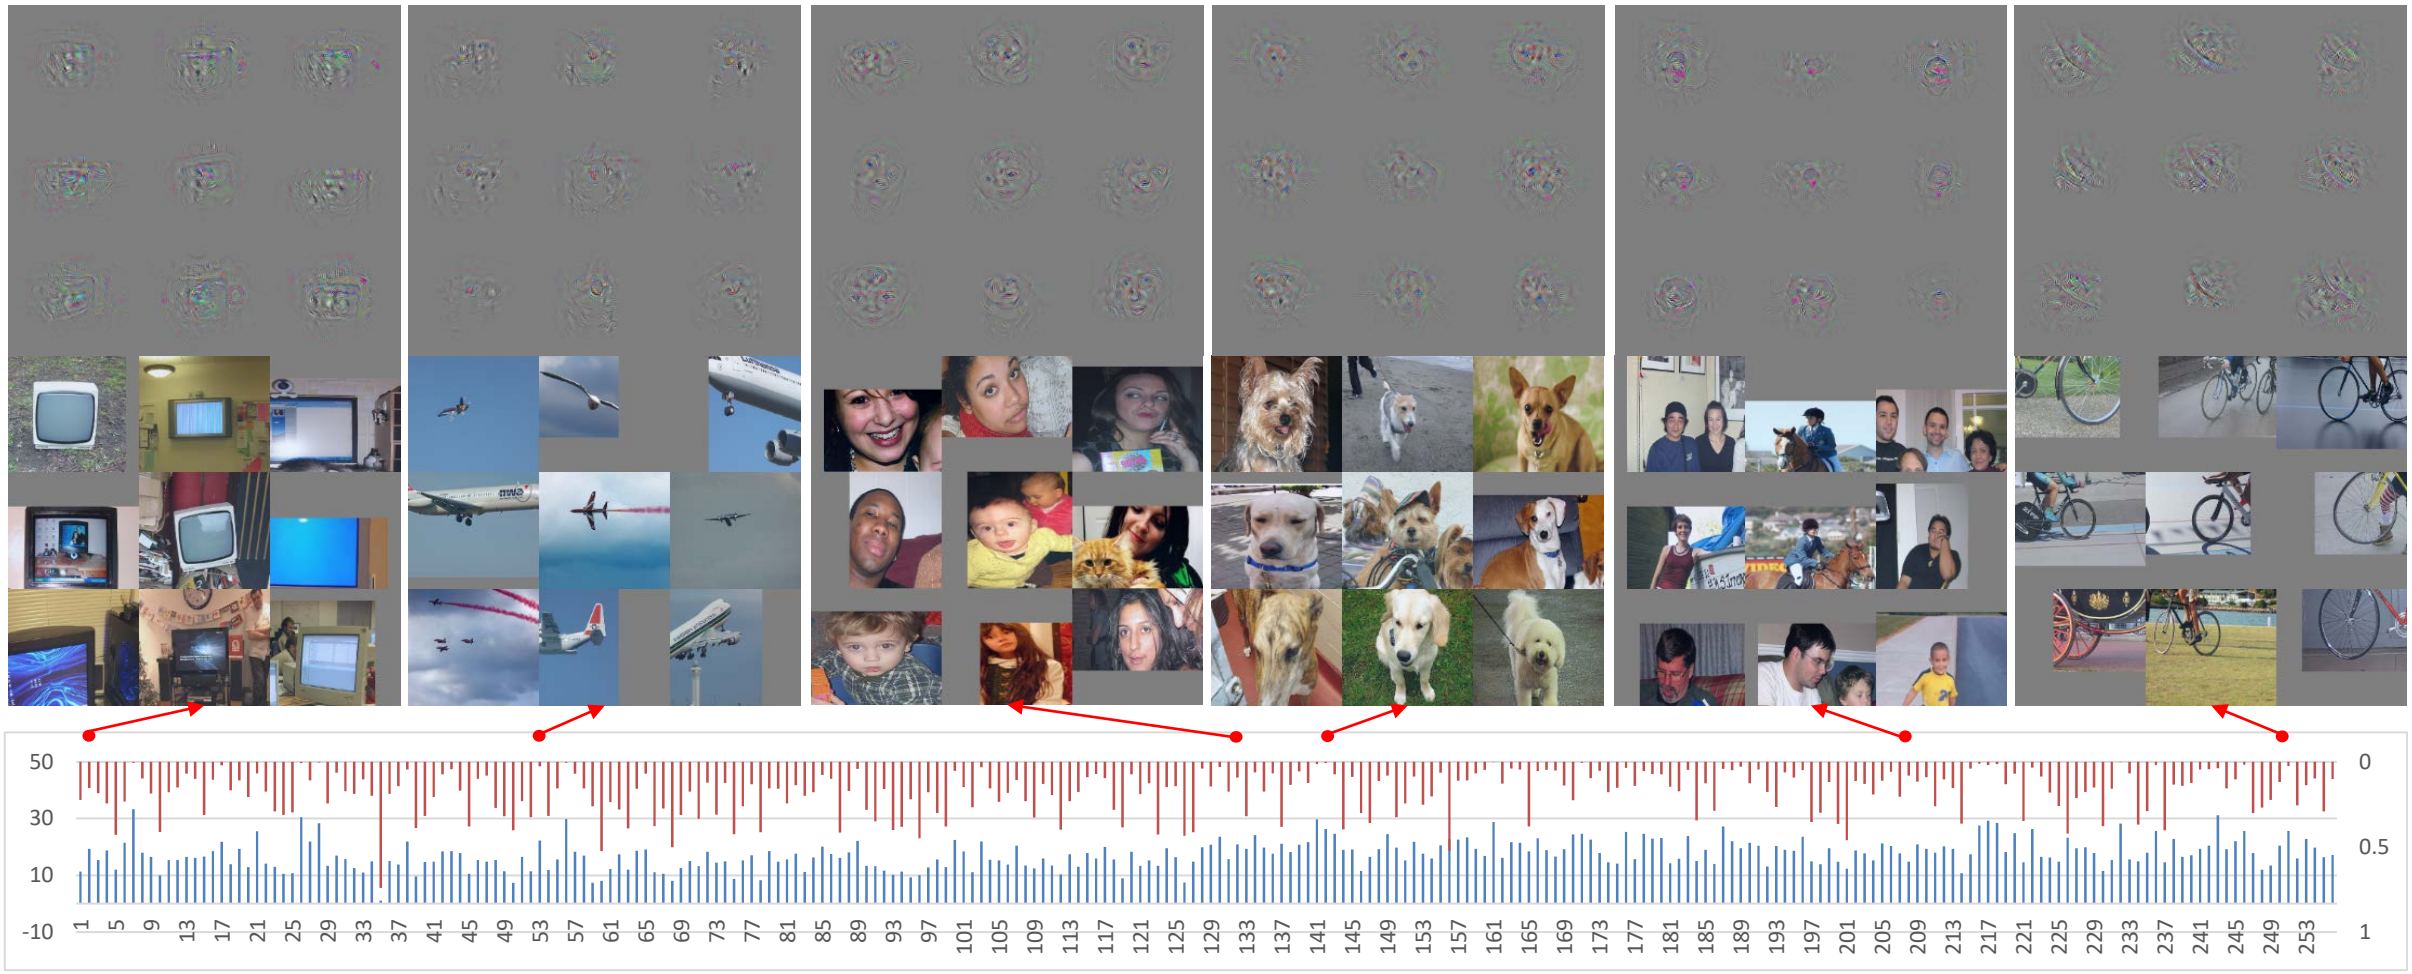

Optimal Conv5 Filters Against Multiple Object Classes (Bad Examples)

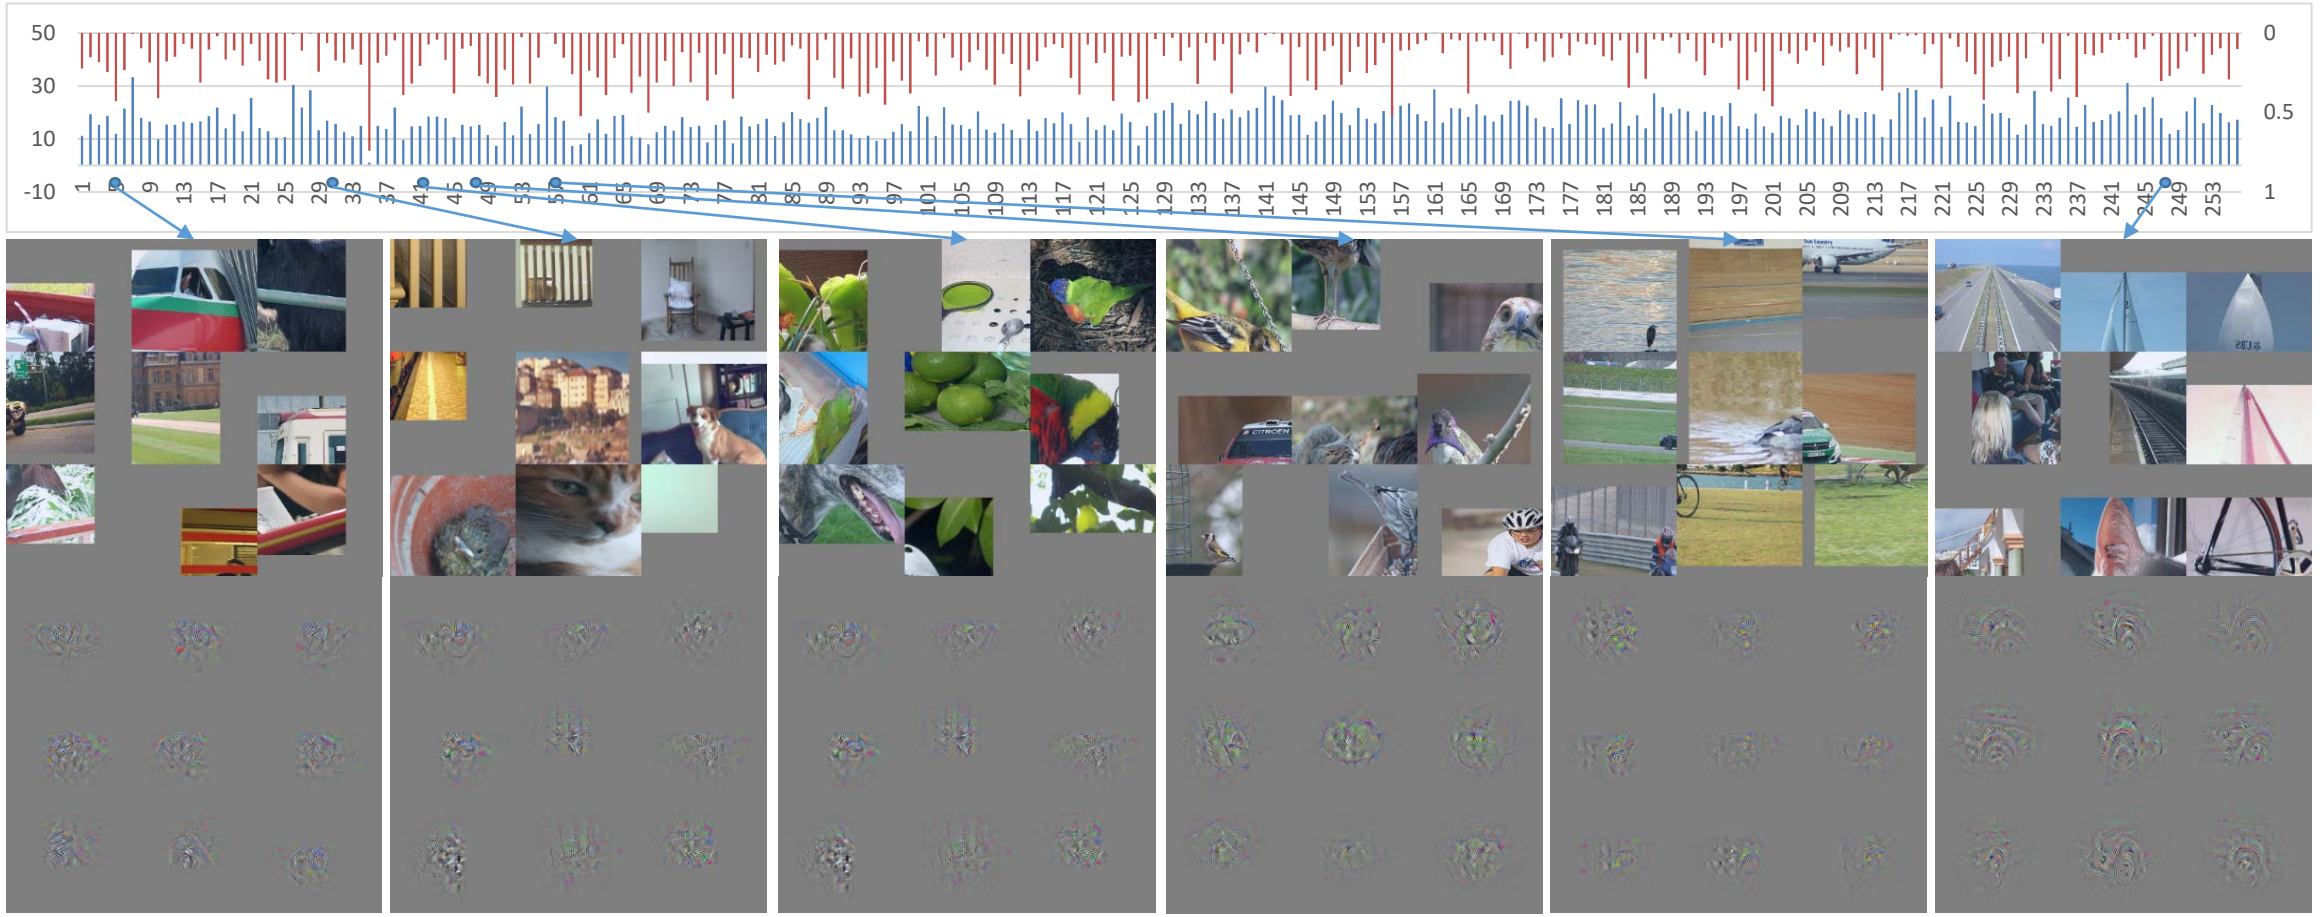

GCM Scores of Top 10 Conv5 Features versus Iteration Number

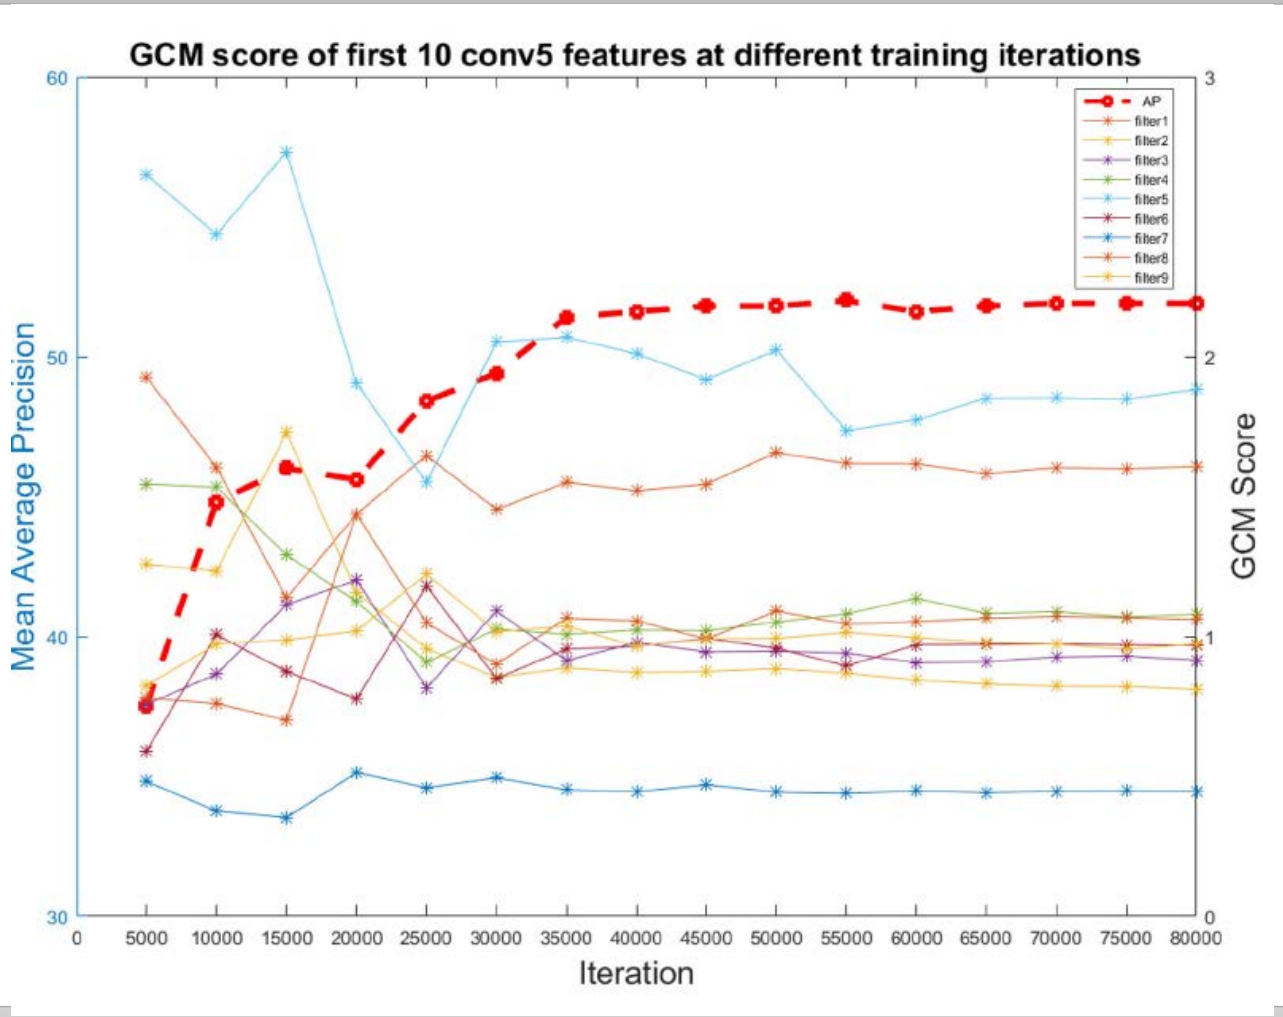

Number of GMC-like Features versus Iteration Number

Network	Iteration	aeroplane	bicycle	bird	boat	bottle	bus	car	cat	chair	cow	table	dog	horse	motorbike	person	plant	sheep	sofa	train	tv	sum
Caffe	10000	4	2	0	1	0	5	0	4	0	2	1	4	4	9	3	1	1	2	1	5	49
	20000	6	3	0	0	0	5	0	8	0	2	2	9	6	10	3	1	1	2	2	5	65
	30000	6	2	0	0	0	6	0	8	0	2	3	5	4	8	3	1	1	2	1	5	57
	40000	6	2	0	1	0	8	2	7	0	2	3	6	5	9	3	2	1	3	3	5	68
	80000	5	2	0	1	0	7	2	9	0	2	5	6	5	10	3	2	1	4	4	5	73
VGG	10000	8	5	0	1	2	11	3	8	0	3	2	6	6	2	4	2	1	3	6	6	79
	20000	6	2	1	1	2	9	3	9	0	2	2	11	3	8	4	1	1	2	4	5	76
	30000	6	3	1	1	1	6	3	4	0	1	2	7	1	1	4	1	1	1	1	4	49
	40000	7	3	1	1	2	9	3	5	0	2	2	6	4	3	3	1	1	1	2	4	60
	80000	8	3	1	1	2	10	3	7	0	3	2	7	5	4	4	2	1	1	2	4	70

Cluster Purity Measure (CPM)

Examples of Cluster Purity Measure (1)

Good

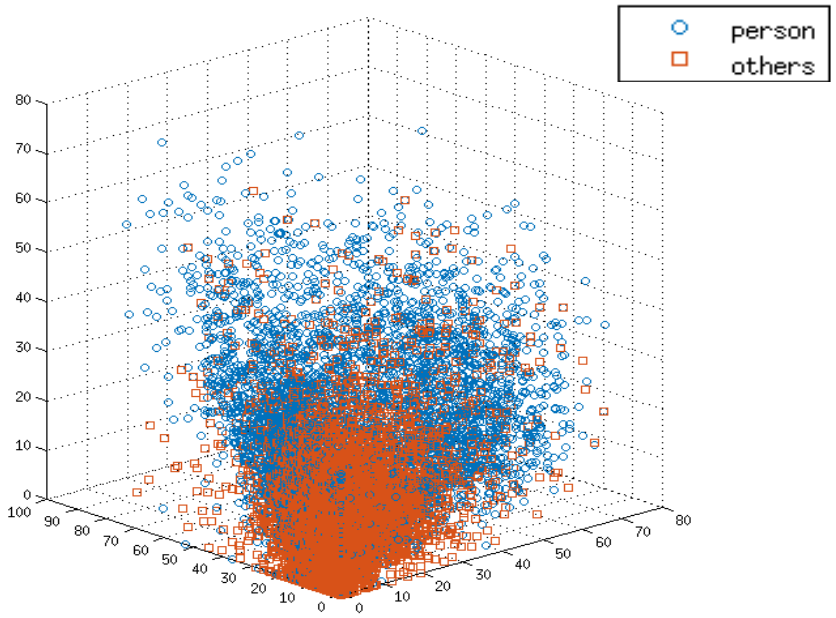

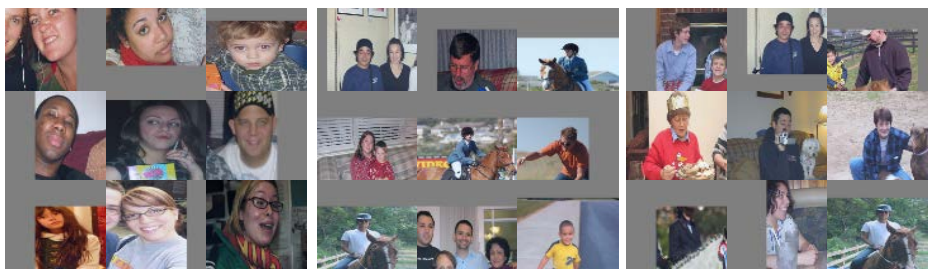

Poor

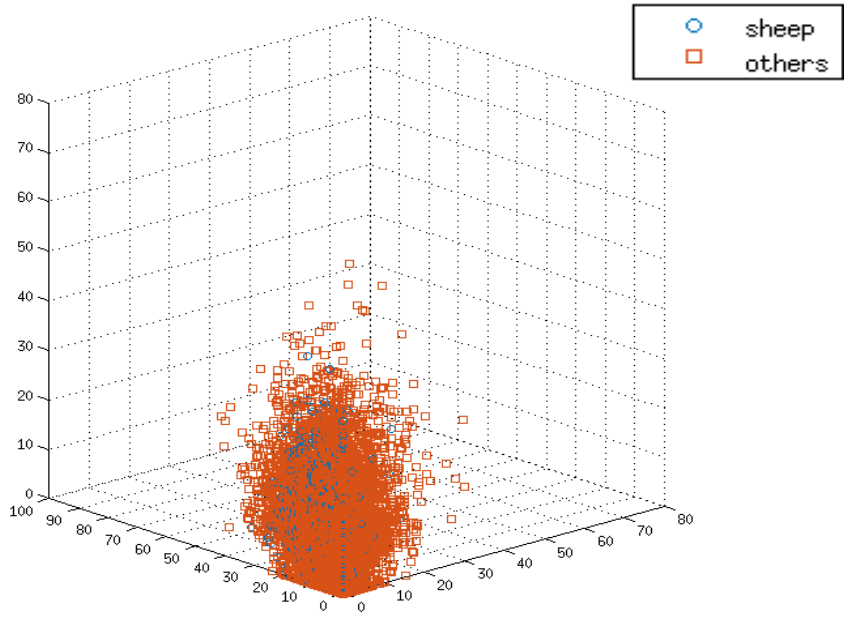

Examples of Cluster Purity Measure (2)

Good

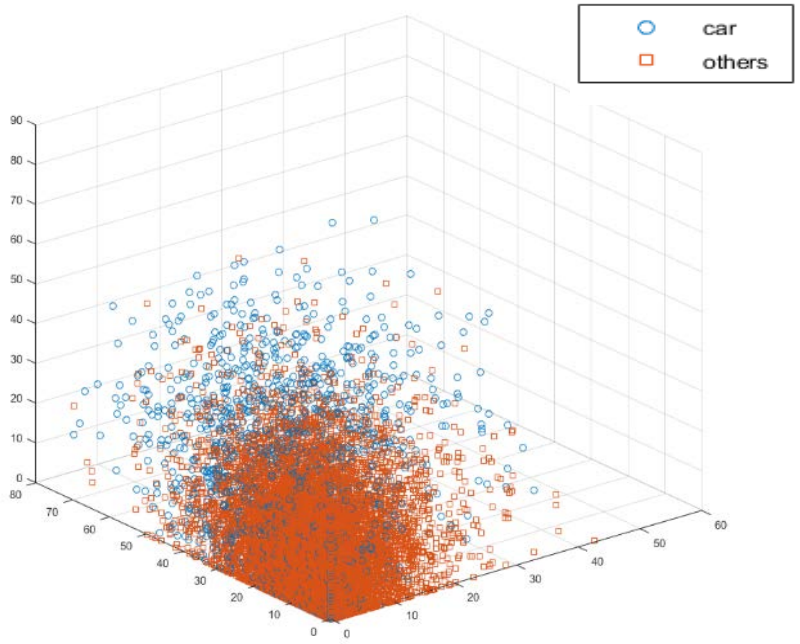

Poor

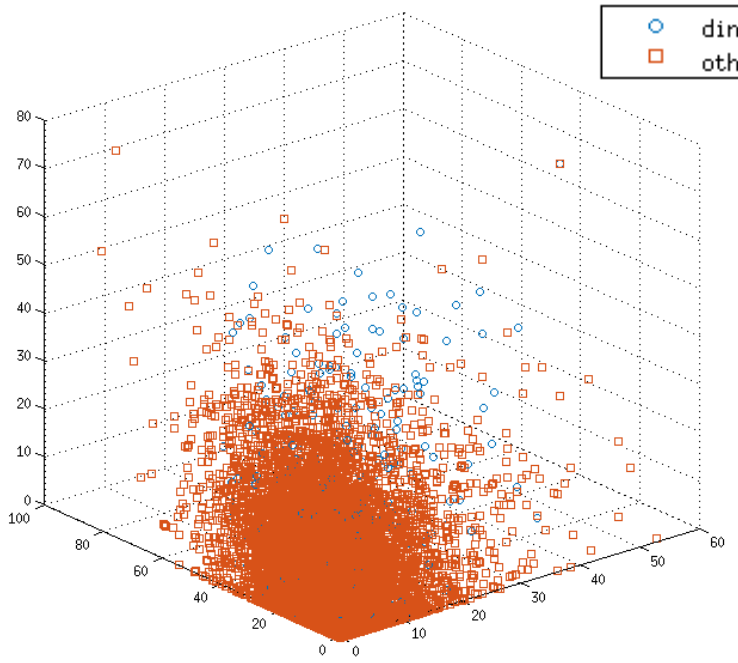

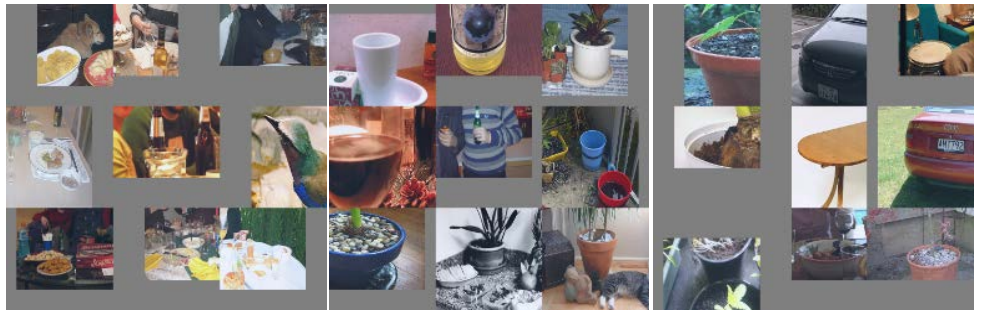

Comments

- Besides feature visualization, we can quantify the power of a particular filter in discriminating an object class
  - 1-D case case: A smaller “Gaussian Confusion Measure (GCM)”
  - Multiple-Dimensional case: A higher “Cluster Purity Measure (CPM)”
- GCM works for a class of objects with a similar setting
- CPM works for a class of objects with multiple settings